

# Security: Applications & Information Flow

Klaus v. Gleissenthall



# Where are we?

- We're in the business of proving things about programs
- First we used Hoare Logic and VCGen
- But this requires supplying loop invariants
- This can be quite tedious!
- Next, we wanted to compute loop invariants (semi-)automatically
- We looked at Horn clauses
- Horn clauses represent constraints on unknown relations called **queries**
- We used weakest preconditions to translate programs into Horn clauses
- In particular, this translation uses queries to represent loop invariants

# Where are we?

- Next, we've looked at solving Horn clauses, that is, compute solution for the unknown queries
- The solutions to the queries give us the missing loop invariants
- We then looked at solving Horn clauses
- First, we start with a solution that maps all queries to  $\perp$
- We then used the **post** operator to compute strongest postconditions
- For each clause, we compute the strongest postcondition of the body wrt. the head
- We can think of this process as computing the set of reachable states
- The strongest postcondition gives us a new state for the solution of the head query
- Unfortunately, if there is an unbounded number of states, the algorithm never terminates

# Where are we?

- To solve this problem we've used abstraction
- Instead of the concrete **post**, we compute an abstract post **post#**
- **post#** expresses the post-condition over a finite vocabulary
- The finite vocabulary is given by a set of predicates **P** (e.g.,  $x \geq 0$ ,  $x \geq 1$ )
- **post#** computes the strongest post-condition that can be expressed as a

conjunction of predicates in **P**

- Restricting the vocabulary ensures that the algorithm always terminates
- It may however be that the computed solution is not strong enough to  
prove our desired post-condition

# Where are we?

- What have we gained?
- Instead of invariants, we now have to supply the building blocks: predicates
- We can just throw a bunch of predicates at it and “see what sticks”
- It doesn’t matter if we supply wrong predicates, as these will just not be used
- However, too many predicates and a large program may make the algorithm slow
- Our algorithm will also guess the Boolean structure of the invariant
- That is, we don’t have to worry about where to put ||

# Where are we?

- Have we solved the problem of computing loop invariants automatically?
  - No! If we're missing the right predicates, the verification will still not work
  - Not always easy to guess, that is, requires some ingenuity
  - Works really well for: large programs, where the invariants have predictable structure
- 
- Does it make verification easier?
  - Don't forget, there can be no perfect solution as we're tackling an undecidable problem.

# Where are we?

- After all, we're proving that for \*any\* possible inputs the program is correct
- These are very strong guarantees, since, as Dijkstra said:

“Program testing can be used to  
show the presence of bugs, but never to show their absence!”

- With verification, we can show their absence, but it's now always easy

# Verification for Security

- Now: How can we use what we've learned for proving security properties?
- First, often proving functional properties alone is good for security!
- Example: verified crypto library <https://dl.acm.org/doi/pdf/10.1145/3133956.3134043>
- Next, we can use reasoning about arrays to prove memory safety
- Memory errors are still among the most exploited and problematic security concerns!
- However, systems code often contains complicated pointer operations  
and difficult data-structures



# Verification for Security

- One direction: pick abstractions that are tailored to memory safety
  - Example: Facebook's Infer <https://fbinfer.com/>
  - Try to detect Null-pointer exceptions in mobile applications
  - Crashes in mobile apps are bad as no one ever updates, so old versions stick around
  - Had some success using this in production
  - Galois: <https://galois.com/> Company that uses formal verification for security & more
  - Amazon/AWS has a huge formal verification group
  - [https://www.youtube.com/watch?v=g-DH\\_b5bFd4](https://www.youtube.com/watch?v=g-DH_b5bFd4)
  - <http://wwwo.cs.ucl.ac.uk/staff/b.cook/oneclick.pdf>
  - Microsoft has the Dafny verifier: <https://github.com/dafny-lang/dafny> (VCGen)
  - SLAM: <https://www.microsoft.com/en-us/research/project/slam/> (Predicate Abstraction)
-

# Applications

- Alternatives for memory safety: Rust (that is, safe subset ;)), Sanitizers
- Current project: a verified sanitizer that provably detects memory errors
- Can be used to correctly patch existing vulnerabilities
- Some more applications:
  - Smart contracts
  - Some devastating hacks!
  - Lots of money on the line

# Applications: Smart Contracts

- Smart contracts are (usually simple, imperative) programs that manage how assets are distributed
- Need to make sure that they don't contain bugs
- Lots of work, and a bunch of start ups:
  - <https://veridise.com/>
  - <https://www.certora.com/>
  - <https://files.sri.inf.ethz.ch/website/papers/sp20-verx.pdf>

# Other Applications

- Proving Security of Crypto Algorithms:
  - <https://github.com/EasyCrypt/easycrypt>
- Verified enclave
  - <https://www.microsoft.com/en-us/research/project/komodo/>
- Operating Systems:
  - <https://sel4.systems/>
- Compiler
  - <https://compcert.org/compcert-C.html>
- Distributed Systems:
  - <https://www.microsoft.com/en-us/research/publication/ironfleet-proving-practical-distributed-systems-correct/>
  - <https://dl.acm.org/doi/10.1145/3290372>

# My Research on Verification

- Verify that a distributed system is correct
- Pretend Synchrony: <https://dl.acm.org/doi/10.1145/3290372>
  - Implement a verified key-value store using Paxos
  - There are other verified implementations, but they require a lot of manual proof
  - Pretend Synchrony exploits insights about algorithms
  - For the purpose of the proof, we can pretend that the algorithm is synchronous
  - Our proofs require a lot less work than previous ones
  - Builds on VCGen, but need some tricks to reason about concurrency

# My Research on Verification

- Iodine: <https://gleissen.github.io/papers/iodine.pdf>
  - Verify that hardware doesn't accidentally leak sensitive information (via timing)
  - Builds on predicate abstraction / Horn clauses
  - But now: the programs we verify are the hardware (written in Verilog!)
  - We can use the same technique by translating into Horn clauses
- Xenon: <https://gleissen.github.io/papers/xenon.pdf>
  - Follow-up that makes the method easier to use and verifies some real RISC-V processors
  - More on this problem now!

# Some more Light & Inspirational Reading:

- <https://www.quantamagazine.org/formal-verification-creates-hacker-proof-code-20160920/>
- <https://www.quantamagazine.org/computing-expert-says-programmers-need-more-math-20220517/>

# Now: Information Flow

- Let's now look at a different set of security problems: side-channels & speculation
  - Can be used to extract information from (otherwise safe) cryptographic functions
  - Real problem in practice, in particular for cloud providers like Amazon/Azure etc.
  - Allow for sneaky attacks that are hard to defend against!
- 
- But let's start with a simpler problem
  - Say, our program contains variables that are **secret**
  - Moreover, there is an attacker, that can **observe** certain other variables
  - We want to make sure that the secrets can't fall into the hands of the attacker



# Information Flow

## Example:

- Consider the following program
- Variable **secret** contains our confidential value
- The attacker can observe variable **obs**

`x := secret;`

`y := x;`

`obs := y`

## Quiz:

- Is this program safe?

# Information Flow

## Example:

- Consider the following program
- Variable **secret** contains our confidential value
- The attacker can observe variable **obs**

`x := secret;`

`y := x;`

`obs := y`

## Quiz:

- Is this program safe?
- No: There is a direct flow from **secret** to **obs** (via variables x, y)

# Information Flow

## Example:

- Let's look at another program
- Again, **secret** contains our confidential value

```
x := secret;  
obs := 0;  
if (x=1) {  
    obs := 1;  
}
```

## Quiz:

- Does this program contain a direct flow?
- Is it safe?
- What can happen?

# Information Flow

## Example:

- Let's look at another program
- Again, **secret** contains our confidential value

```
x := secret;  
obs := 0;  
if (x=1) {  
    obs := 1;  
}
```

- The program contains an indirect flow
- Depending on whether we take the branch or not, **obs** will have value 0 or 1.
- This leaks information about **obs**: the attacker can find out if **secret** has value 1

# Non-Interference

- When is a program safe, according to this notion?
- We can formalize this in the notion of non-interference
- Intuition: let's look back at our initial unsafe example
- Let's look at two executions of the program (left and right)
- In these two executions the **secrets** take different values
- All other values are the same
- The program is safe, if the observables have the same value
- That is, despite the **secrets** being different, an attacker cannot distinguish the two runs

$x := \text{secret};$

$y := x;$

**obs**  $:= y$

$x := \text{secret};$

$y := x;$

**obs**  $:= y$

# Non-Interference

## Example:

$x := 3;$

$y := x;$

$\text{obs} := y$

$x := 5;$

$y := x;$

$\text{obs} := y$

- Say we set the secret to **3** in the left run and **5** in the right run
- After executing the program **obs** is **3** in the left run and **5** in the right run
- Hence, an attacker can distinguish the two runs, and the program is not safe!

# Non-Interference

Example:

```
x := secret;
```

```
obs := 0;
```

```
if (x=1) {
```

```
    obs := 1;
```

```
}
```

```
x := secret;
```

```
obs := 0;
```

```
if (x=1) {
```

```
    obs := 1;
```

```
}
```

**Quiz:**

- Let's look at our second example
- Is the example safe according to our new notion of security?

# Non-Interference

Example:

```
x := 1;  
obs := 0;  
if (x=1) {  
    obs := 1;  
}
```

```
x := 5;  
obs := 0;  
if (x=1) {  
    obs := 1;  
}
```

- No!
- If we pick secret **1** in the left run and **5** in the right run, the observations will differ
- In the left run, the attacker observes **1**, in the right run **0**
- This is called an indirect flow
- Note, we have to pick specific values for the difference to occur
- For the program to satisfy non-interference attacker observations need to be the same **for any** choice of secrets!



# Non-Interference

## Quiz:

- Remember our semantics relation  $\langle s, \sigma \rangle \Downarrow \sigma'$ , saying state  $\sigma$  changes to state  $\sigma'$  after executing statement  $s$
- Let's write  $\sigma \downarrow \text{obs}$  for restriction of  $\sigma$ 's domain to attacker observable variables
- How can we formally define non-interference?

# Non-Interference

## Quiz:

- Remember our semantics relation  $\langle s, \sigma \rangle \Downarrow \sigma'$ , saying state  $\sigma$  changes to state  $\sigma'$  after executing statement  $s$
- Let's write  $\sigma \downarrow \text{obs}$  for restriction of  $\sigma$ 's domain to attacker observable variables
- How can we formally define non-interference?

For all  $\sigma_L, \sigma_R$ : if  $\sigma_L \downarrow \text{obs} = \sigma_R \downarrow \text{obs}$  and  $\langle s, \sigma_L \rangle \Downarrow \sigma_L'$  and  $\langle s, \sigma_R \rangle \Downarrow \sigma_R'$  then  $\sigma_L' \downarrow \text{obs} = \sigma_R' \downarrow \text{obs}$

# Non-Interference

Quiz:

**obs** := **obs** + 1;

- Does the program above satisfy non-interference (as we defined it)?

# Non-Interference

## Quiz:

```
x := secret;  
while (x ≥ 1) {  
    x := x + 1;  
}
```

- Does the program above satisfy non-interference (as we defined it)?
- Could an attacker still distinguish two runs with different secrets?

# Non-Interference

## Quiz:

```
x := secret;  
while (x ≥ 1) {  
    x := x + 1;  
}
```

- Does the program above satisfy non-interference (as we defined it)?
- Could an attacker still distinguish two runs with different secrets?
- Yes! If an attacker can see whether the program terminates, they can distinguish the runs
- We defined termination insensitive non-interference
- One can define termination sensitive non-interference, which in addition requires the program to terminate

# Self-composition

## Quiz:

$x := \text{secret};$

$y := x;$

$\text{obs} := y$

$x := \text{secret};$

$y := x;$

$\text{obs} := y$

- Let's look at our example program
- Can we use our previous verification techniques to verify non-interference?

# Self-composition

## Quiz:

$x := \text{secret};$

$y := x;$

$\text{obs} := y$

$x := \text{secret};$

$y := x;$

$\text{obs} := y$

- Let's look at our example program
- Can we use our previous verification techniques to verify non-interference?
- Not directly! Non-interference talks about two runs
- This is called a hyper-property
- But, we can form a new program made up of two copies
- We can then use verification to prove non-interference
- This is called self-composition or product program

<https://www.sop.inria.fr/lemme/Tamara.Rezk/publication/Barthe-DArgenio-Rezk-Journal.pdf>

# Information Flow Types

- We'll now take a look at an approach with a different philosophy
- Up to now, you can write *any* program and we'll try to prove correctness
- This is the verification approach (conferences like CAV)
- There's a different approach that comes from the PL community (PLDI, POPL)
- You're not allowed to write programs that might be wrong
- Idea: you cannot even express possibly incorrect programs
- Comes at the risk of ruling out programs that are correct but less obviously so
- Totalitarian programming: wrong programs are thought crime
- Another example: Rust
- For security, this is called language-based security



“In the end we shall make thoughtcrime literally impossible, because there will be no words in which to express it.”

**George Orwell (1984)**



# Information Flow Types

- We'll now define a type system such that if the program type checks, then it is non-interferent
- We'll do this for Nano:

**Expressions:**  $e, e_1, e_2 \ni \text{Exp} ::= n \mid x \mid e_1 + e_2 \mid e_1 - e_2 \mid e_1 * e_2$  where  $n \in \mathbb{Z}, x \in \text{Vars}$

**Boolean Expressions:**  $b, b_1, b_2 \ni \text{BExp} ::= \top \mid \perp \mid \neg b \mid b_1 \wedge b_2 \mid b_1 \vee b_2 \mid e_1 = e_2 \mid e_1 \leq e_2$

**Statement:**  $s, s_1, s_2 \ni \text{Stmt} ::=$

skip	(no-op)
$x := e$	(assignment)
$s_1 ; s_2$	(sequential composition)
if $b$ then $s_1$ else $s_2$	(if)
while $b$ do $s$	(while)

# Information Flow Types

- First, we'll note that our two levels of security **sec** and **obs** form a lattice  $\mathcal{L}$
- That means we can define a relation  $\sqsubseteq$  such that  $a \sqsubseteq b$ , iff  $b$  is at least as confidential as  $a$ .
- You can think of  $b$  as requiring higher (or equal) “security clearance”

## Quiz:

- How should  $\sqsubseteq$  be defined for **sec** and **obs** ?

- We also define a join operator  $\sqcup$  such that for all  $a, b \in \mathcal{L}$ :  $a \sqsubseteq a \sqcup b$  and  $b \sqsubseteq a \sqcup b$

## Quiz:

- How should  $\sqcup$  be defined for **sec** and **obs** ?

# Information Flow Types

- Let's assume we have a function  $\Gamma$  that assigns a security label to each variable

Quiz:

$x := \text{secret};$

$y := x;$

$\text{obs} := 3$

- Let's look again at the following example:
- What should we assign to  $\Gamma(\text{secret})$ ?
- What about  $\Gamma(\text{obs})$ ?
- What about  $\Gamma(x)$ ?
- What about  $\Gamma(y)$ ?
- Is the program safe?

# Information Flow Types

- Let's assume we have a function  $\Gamma$  that assigns a security label to each variable

Quiz:

$x := \text{secret};$

$y := x;$

$\text{obs} := y$

- What if, we change it to match our previous example
- What's the type of  $\Gamma(y)$  now?
- Is the program safe?
- Now: formalize this in type system

# Information Flow Types

- Let's define a typing judgement  $\Gamma \vdash e : \ell$  saying that expression  $e$  has label  $\ell \in \mathcal{L}$
- We define them using inference rules, one per expression

$$\begin{array}{c}
 \text{var} \quad \frac{\Gamma(x) = \ell}{\Gamma \vdash x : \ell} \qquad \text{num} \quad \frac{}{\Gamma \vdash n : \text{obs}} \qquad \text{plus} \quad \frac{\Gamma \vdash e_1 : \ell_1 \quad \Gamma \vdash e_2 : \ell_2}{\Gamma \vdash e_1 + e_2 : \ell_1 \sqcup \ell_2}
 \end{array}$$

- Similarly, for binary a expression  $b$ , we get

$$\begin{array}{c}
 \text{I} \quad \frac{}{\Gamma \vdash \top : \text{obs}} \qquad \text{or} \quad \frac{\Gamma \vdash b_1 : \ell_1 \quad \Gamma \vdash b_2 : \ell_2}{\Gamma \vdash b_1 \vee b_2 : \ell_1 \sqcup \ell_2} \qquad \text{eq} \quad \frac{\Gamma \vdash e_1 : \ell_1 \quad \Gamma \vdash e_2 : \ell_2}{\Gamma \vdash e_1 = e_2 : \ell_1 \sqcup \ell_2}
 \end{array}$$

# Information Flow Types

## Quiz:

- Let's say  $\Gamma(x) = \text{obs}$  and  $\Gamma(y) = \text{sec}$
- Can we derive that  $\Gamma \vdash x : \text{obs}$  ?
- Can we derive that  $\Gamma \vdash x : \text{sec}$  ?
- What label  $\ell$  do we get for  $\Gamma \vdash 3 : \ell$  ?
- What label  $\ell$  do we get for  $\Gamma \vdash x + 3 : \ell$  ?
- What label  $\ell$  do we get for  $\Gamma \vdash y + 3 : \ell$  ?
- What label  $\ell$  do we get for  $\Gamma \vdash x + y : \ell$  ?

# Information Flow Types

- All for today
- Next lecture: typing statements
- First a wrong attempt, then the correct one
- Then, how to use type systems to rule out side-channel (and speculation) attacks

# Reading

- Lattices: <https://ieeexplore.ieee.org/document/246638>
- IFC Type systems: <https://www.cs.cornell.edu/andru/papers/jsac/sm-jsac03.pdf>
- Lecture notes: [https://www.cs.uoregon.edu/research/summerschool/summer19/lecture\\_notes/myers1.pdf](https://www.cs.uoregon.edu/research/summerschool/summer19/lecture_notes/myers1.pdf)
- There are also videos available for these lectures:
- <https://www.cs.uoregon.edu/research/summerschool/summer12/curriculum.html>
- Under: Language-based security